

Lecture 2: Backgrounds & Preliminaries

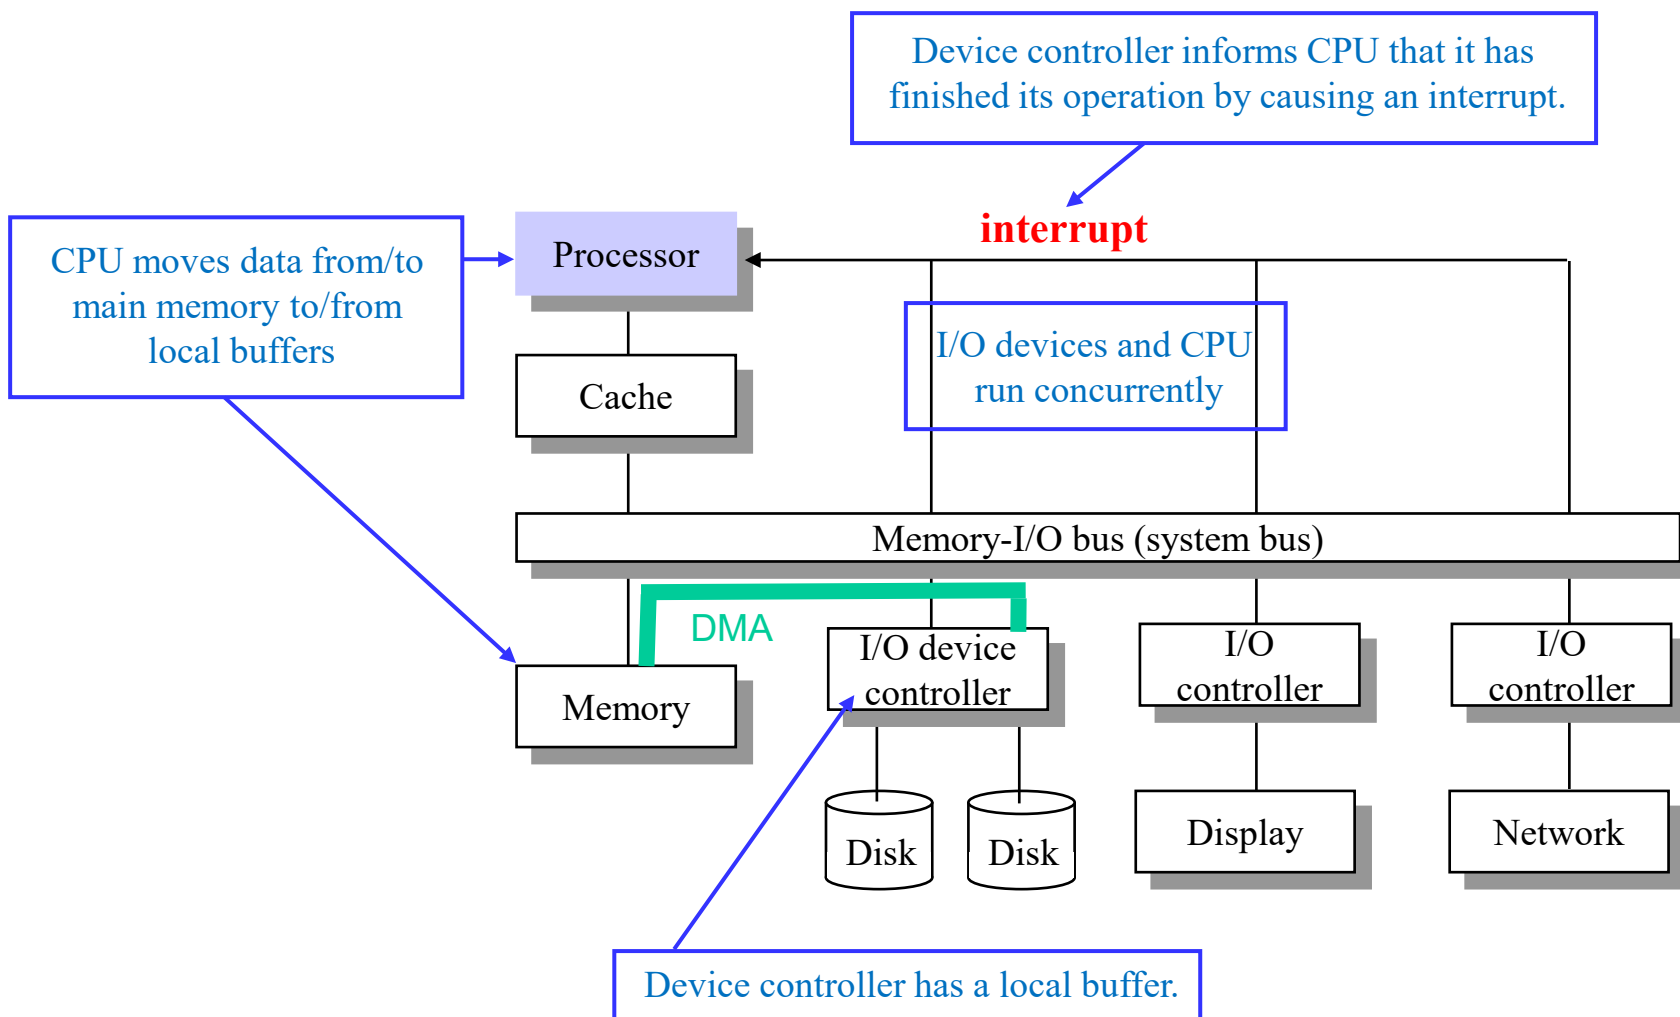
Spring 2020

차호정

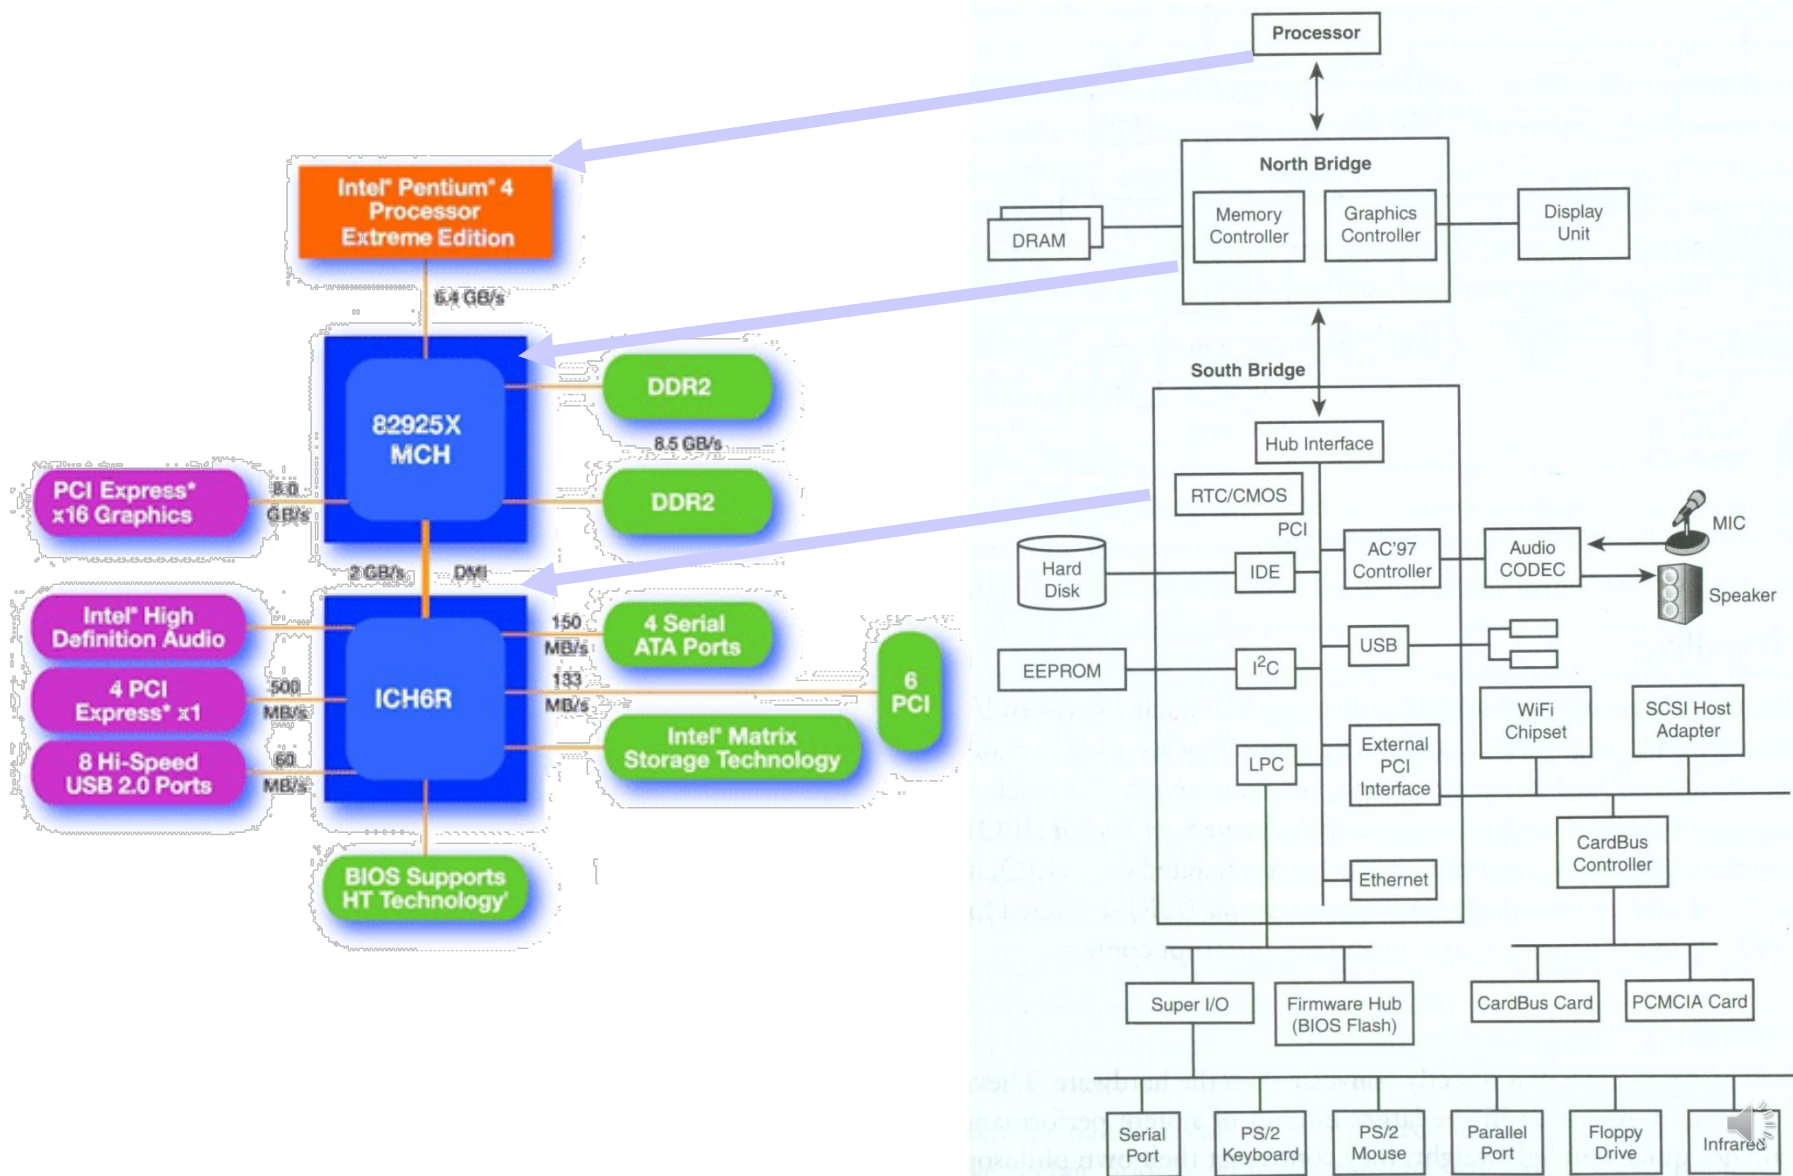
연세대학교 컴퓨터과학과




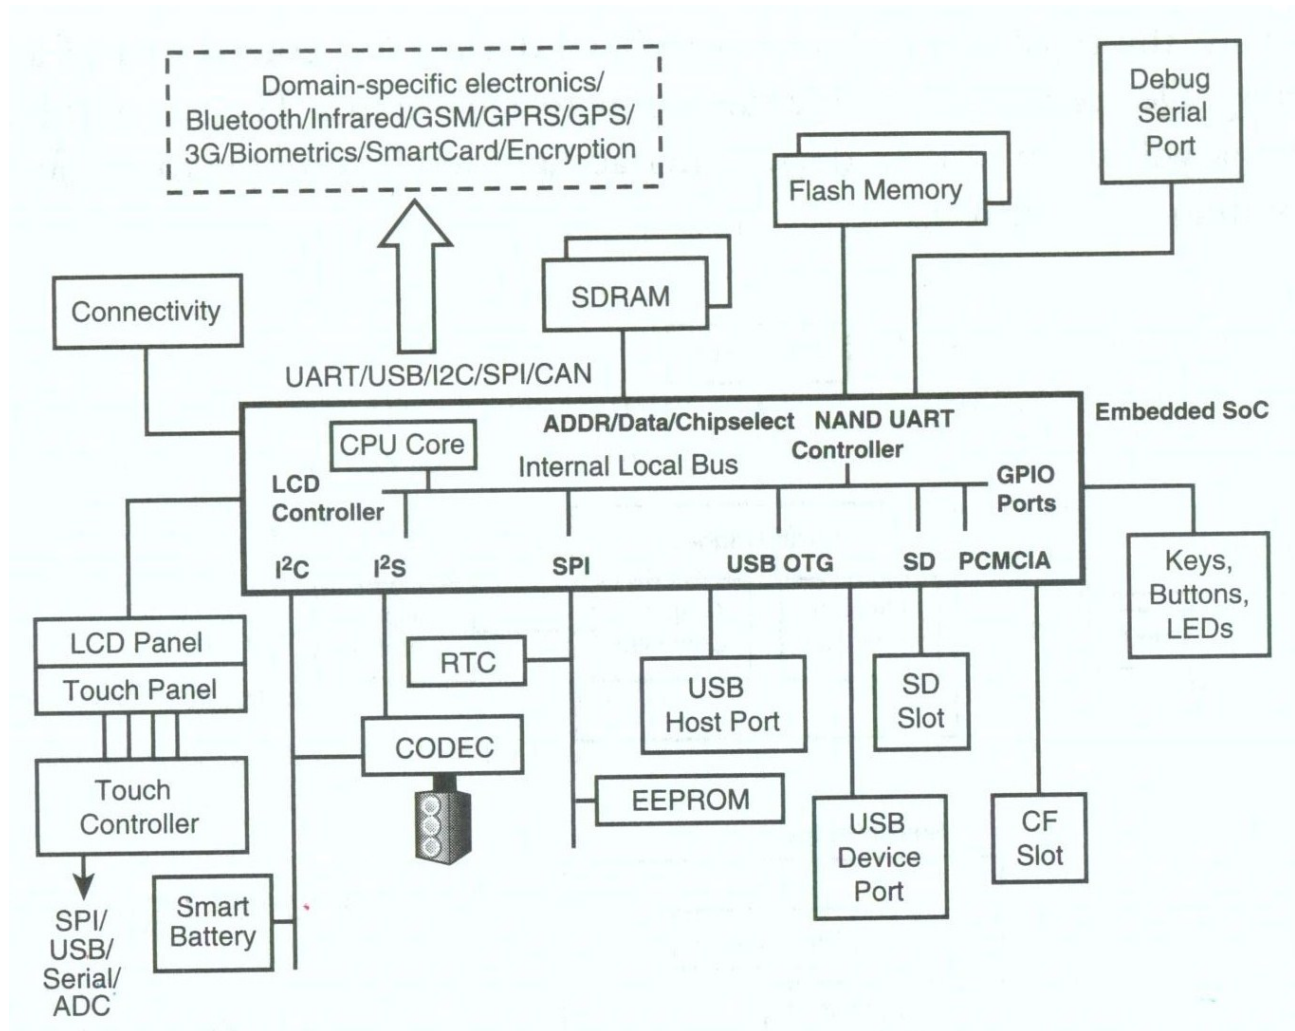
Computer System Architecture: Conventional View



Real World: a Large Pentium (PC) System

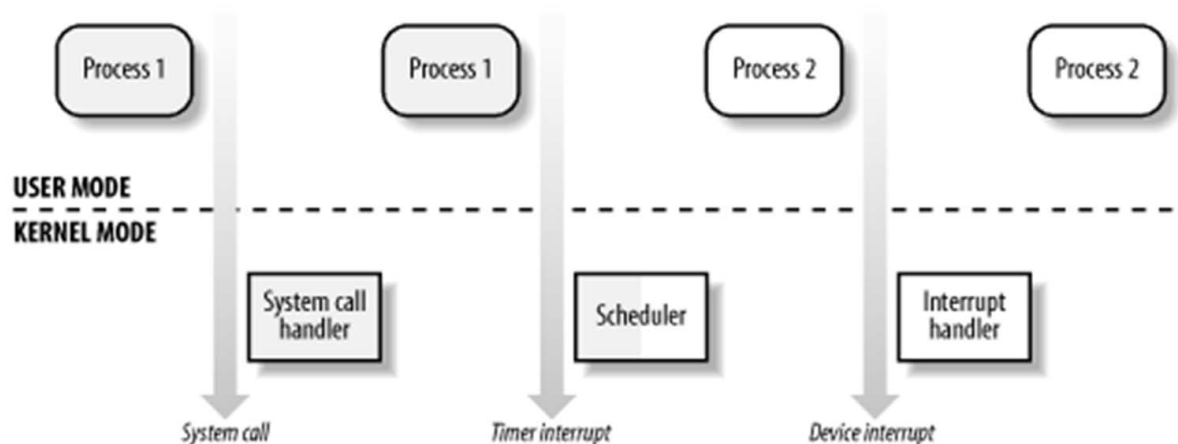


Real World: Embedded System



Hardware Protection for OS

- ‘Dual mode’ operation
 - A properly designed OS must ensure that an incorrect program cannot cause other programs to execute incorrectly.
 - OS has a dual mode mechanism (*user mode* and *kernel mode*) to protect the system from many types of application faults
 - Processor provides the protection mechanism



Dual-Mode (User/Kernel Mode) Operation (1)

- Provide *hardware support* to differentiate between at least two modes of operations
 - **User mode**: execution done on behalf of a user
 - **Kernel mode** (also “*supervisor mode*” or “*system mode*” or “*monitor mode*”): execution done on behalf of operating system
 - Mode is set by a status bit in a protected processor register.
 - Example:
 - Intel 80x86 has four different execution states (i.e. the `cs` segmentation register includes 2-bit field that specifies the CPL (Current Privilege Level) of the CPU.
 - All standard UNIX kernels make use of only User Mode and Kernel Mode.
 - Some machine instructions are designed as *privileged (protected) instructions* and they can be issued only in kernel mode.
 - Most of the I/O instructions

Dual-Mode (User/Kernel Mode) Operation (2)

- Crossing protection boundaries
 - User programs must call on OS to do something privileged (i.e., invoking privileged instructions)
 - Pass control to a kernel service routine running in kernel mode.
 - The kernel verifies that the parameters are correct and legal, executes the request.
- Three cases: *(talk about this in the coming lectures)*
 - (1) Hardware Interrupt
 - (2) Software interrupt (exception)
 - (3) System call

Case Study: x86 *Real* Mode

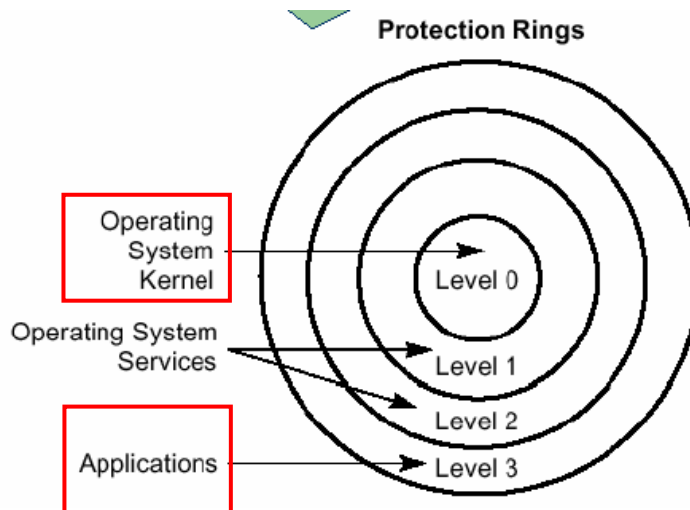
- Real mode refers to compatibility with 16 bit Intel CPUs (8086, 80286)
- All x86 CPUs start in “real mode”
 - The system BIOS only works in real mode.
 - So, the boot code has to work in real mode.
- Segmented memory
 - All segments are restricted to 64KB in size when in real mode
 - IP (instruction pointer), segment registers are all 16 bits!
 - Awkward to work with objects larger than a segment (64KB)
 - No paging or memory protection

Case Study: x86 *Protected* Mode

- The 386 and higher CPUs have a *protected mode*
 - 32bit memory address
 - Memory protection
 - Virtual memory paging
 - IO protection
 - Privilege levels
 - Task switching
 - Interrupt handling
- Entering protected mode
 - Construct valid code, data, and stack segments (GDT, LDT)
 - Set PE bit in CR0 register
 - Jump to a valid code address in a code segment

x86 Protected Mode: Privilege Level

- Privilege checking
 - Ensure that the currently-executing program cannot access areas of memory unless permitted to do so.



- 0 is most privileged, 3 is least
- Most OS uses 0 for kernel code, 3 for user code
- Privilege levels 1 and 2 could be used for more fine-grained protection (e.g., device drivers)

- Three components are involved in the privilege checking:
 - CPL (current privilege level) of the current program
 - RPL (requestor privilege level) in the segment register
 - DPL (descriptor privilege level) of the target segment

Operating System Services (1)

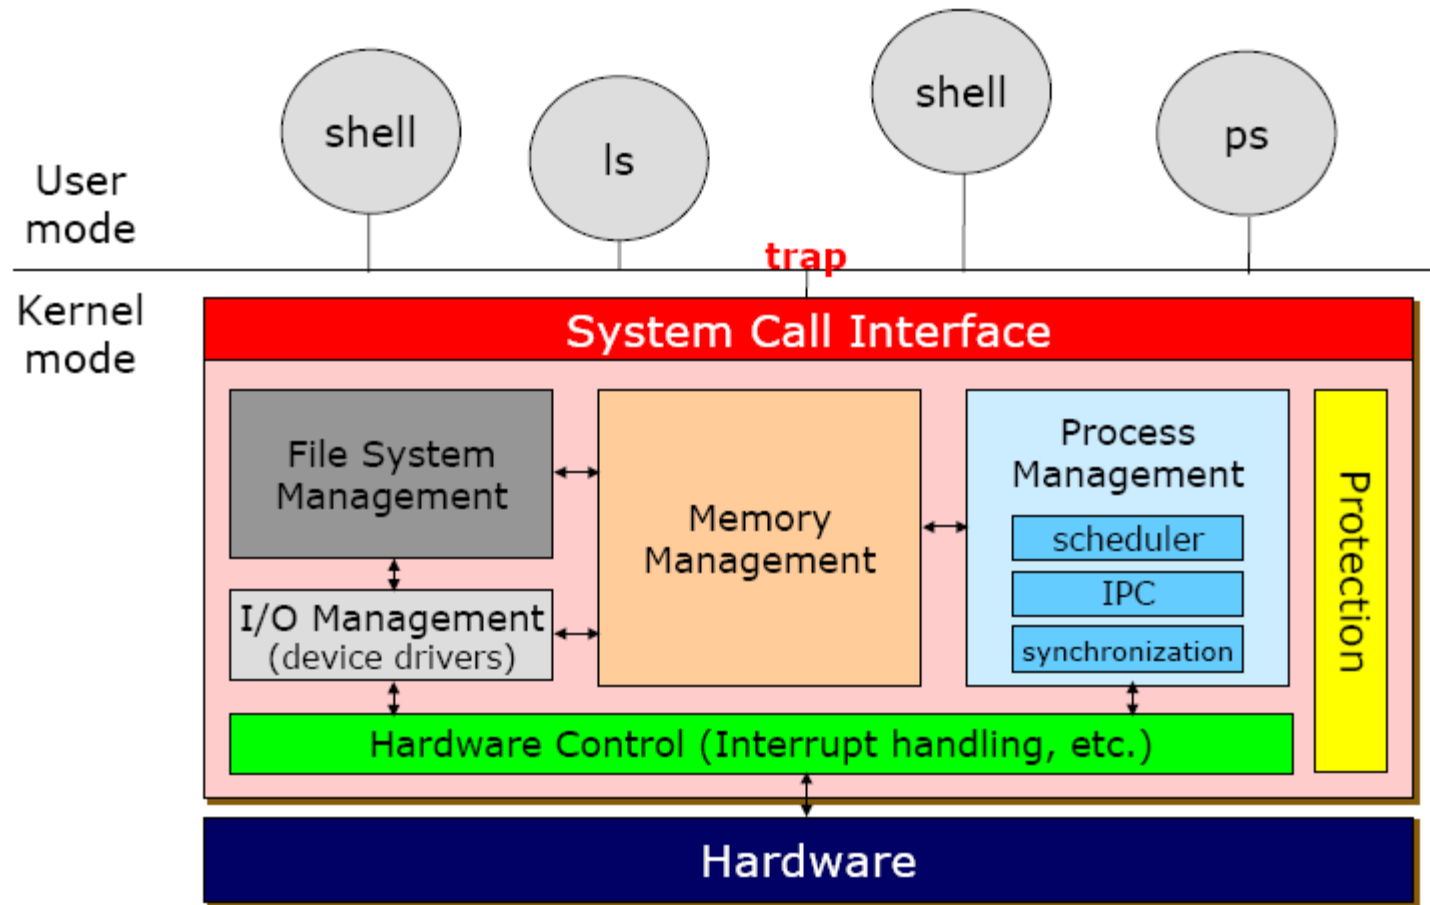
- User Services

- Program execution
 - System capability to load a program into memory and run it.
- I/O operations
 - Since user programs cannot execute I/O operations directly, the operating system must provide some means to perform I/O.
- File-system manipulation
 - Program capability to read, write, create, and delete files.
- Communications
 - Exchange of information between processes executing either on the same computer or on different systems tied together by a network.
- Error detection
 - Ensure correct computing by detecting errors in the CPU and memory hardware, in I/O devices, or in user programs.

Operating System Services (2)

- Resource allocation
 - Allocating resources to multiple users or multiple jobs running at the same time.
- Accounting
 - Keep track of and record which users use how much and what kinds of computer resources for account billing or for accumulating usage statistics.
- Protection
 - Ensuring that all access to system resources is controlled

Operating System Structure



System Call Interface

- ‘System Calls’
 - Provide the interface between a running program and the operating system.
 - Generally available as function calls.

Process Management	fork	CreateProcess	Create a new process
	waitpid	WaitForSingleObject	Wait for a process to exit
	execve	(none)	CreateProcess = fork + execve
	exit	ExitProcess	Terminate execution
	kill	(none)	Send a signal
File Management	open	CreateFile	Create a file or open an existing file
	close	CloseHandle	Close a file
	read	ReadFile	Read data from a file
	write	WriteFile	Write data to a file
	lseek	SetFilePointer	Move the file pointer
	stat	GetFileAttributesEx	Get various file attributes
	chmod	(none)	Change the file access permission
File System Management	mkdir	CreateDirectory	Create a new directory
	rmdir	RemoveDirectory	Remove an empty directory
	link	(none)	Make a link to a file
	unlink	DeleteFile	Destroy an existing file
	mount	(none)	Mount a file system
	umount	(none)	Unmount a file system
	chdir	SetCurrentDirectory	Change the current working directory

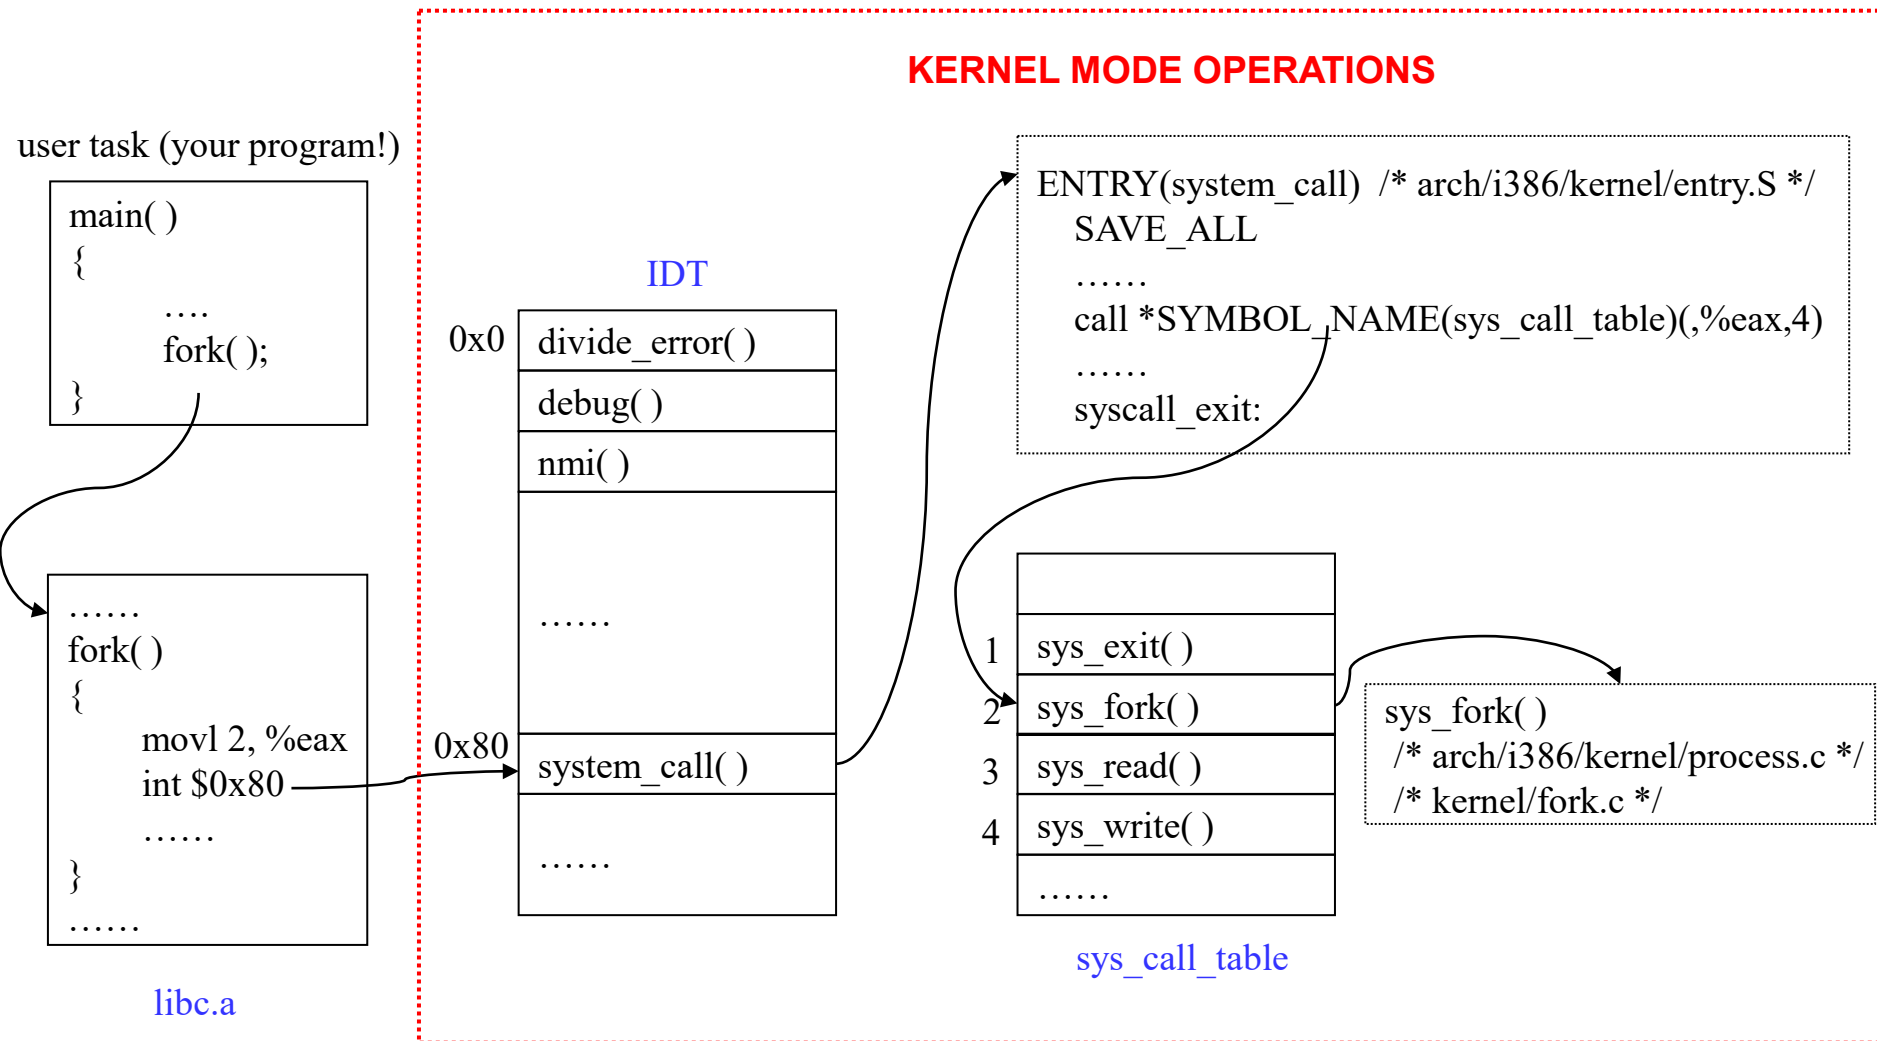
System Call Principles

- Putting an extra layer between the applications and hardware
 - Advantages
 - **Easy to program**: freeing user from aware low-level programming characteristics of hardware devices
 - **Increasing system security**: the kernel can check the correctness of the request at the interface level
 - **Increase program portability**
- System calls
 - UNIX systems implement most interfaces between User Mode processes and hardware devices by means of **system calls** issued to the kernel.
 - Interfaces between User Mode processes and hardware devices
 - To request the kernel services

POSIX APIs and System Calls

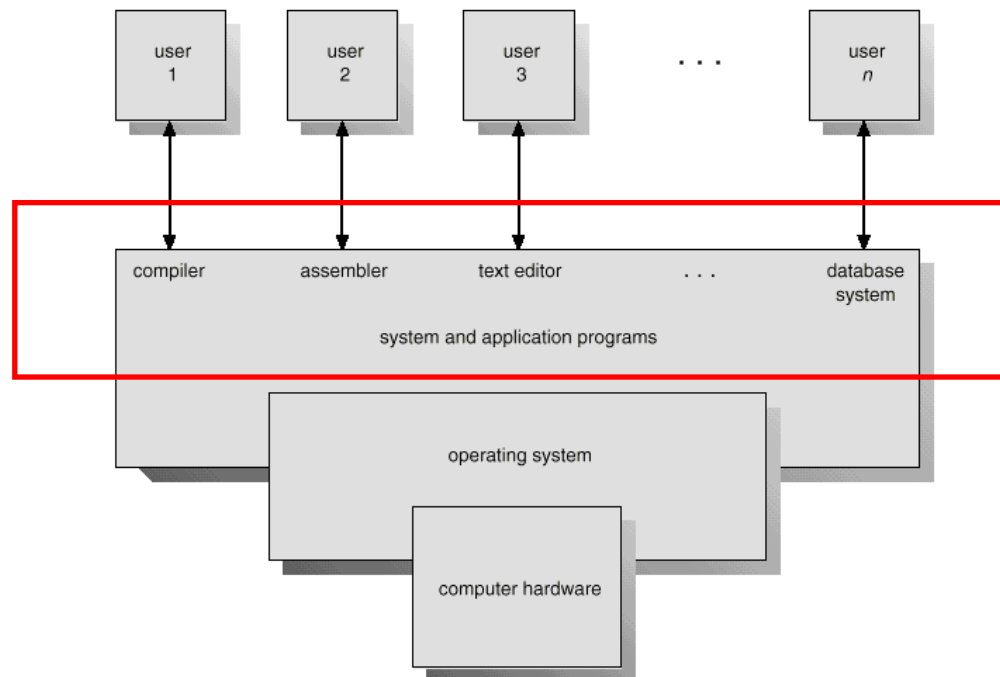
- API (Application Programming Interface)
 - A function definition that specifies how to obtain a given service
 - For example, POSIX APIs - `malloc()`, `calloc()`, `free()` - are implemented in the *libc*, which uses the `brk()` system calls.
 - “POSIX-compliant” if a system offers the proper set of APIs to the applications, no matter how the corresponding functions are implemented.
 - Programmers point of view: User Mode libraries
- System call
 - An explicit request to the kernel made via a software interrupt.
 - Kernel designer's pointer of view: belongs to the kernel
 - Some system calls takes one or more arguments.
 - Returns an integer value
 - If failed: return `-1` and set `errno` (see `include/asm-i386/errno.h`)
 - Implementation: a system call is implemented as a wrapper function in *libc* in user space.

System Call Handling in x86/Linux



OS & System Programs (1)

- ‘System Programs’
 - Provide a convenient environment for program development and execution.
 - Most users’ view of the operating system is defined by system programs, not the actual system calls.



OS & System Programs (2)

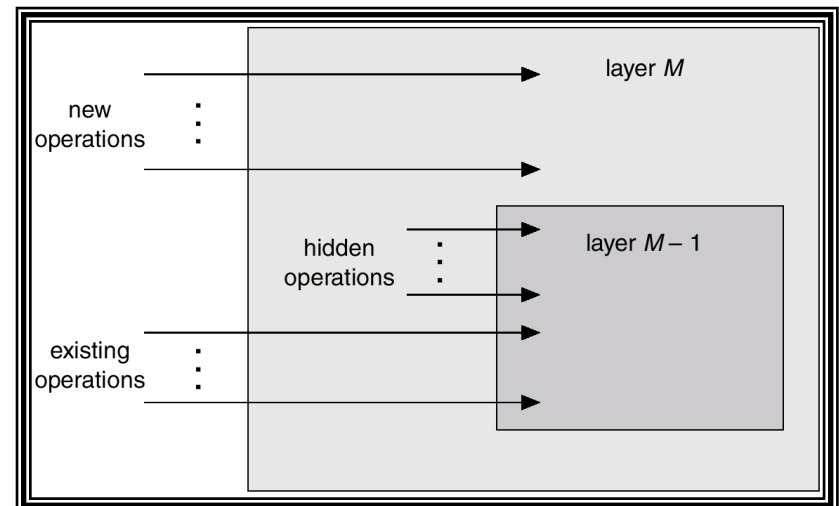
- System programs can be classified into:
 - File manipulation (i.e., shell commands in UNIX)
 - create, delete, copy, rename, print, dump, list, ...
 - Status information
 - date, df, du, top, ...
 - File modification
 - Editor
 - Programming language support
 - Compiler, assembler, interpreter
 - Program loading and execution
 - Loader, linkage editor, debugger
 - Application programs
 - DBMS, Web browsers, word processors, games, plotting tools, ...

Operating System Design (1)

- System Design Goals
 - User Goals
 - Operating system should be convenient to use, easy to learn, reliable, safe, and fast.
 - System Goals
 - Operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient.
- Mechanisms and Policies
 - *Mechanisms* determine how to do something, *policies* decide what will be done.
 - Example
 - Mechanism: Time slicing in time-sharing system
 - Policy: time quantum
 - The separation of policy from mechanism is an important principle.
 - Allow maximum flexibility if policy decisions are to be changed later.

Operating System Design (2)

- Layering Approach
 - The operating system is divided into a number of layers (levels), each built on top of lower layers.
 - Advantages
 - Modularity: layers are selected such that each uses functions (operations) and services of only lower-level layers.
 - Simplify debugging and system verification
 - Difficulties
 - Careful definition of layers
 - Performance



Operating System Design (3)

- System Implementation
 - Traditionally written in assembly language, operating systems can now be written in higher-level languages.
- Code written in a high-level language:
 - Can be written faster.
 - Easier to understand and debug.
 - Far easier to *port* !

Microkernels (1)

- Monolithic Kernel

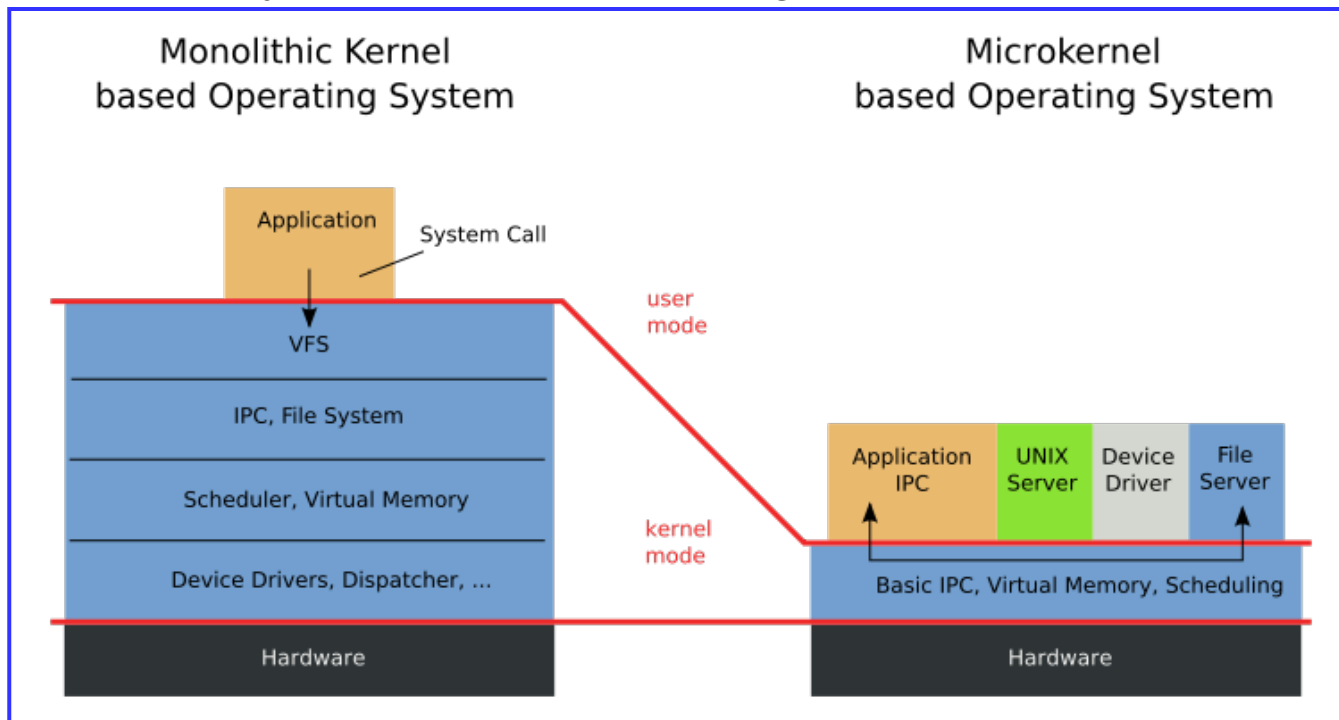
- All operating system services implemented in one big monolithic kernel
- Virtually any procedure calls any other procedure.

- Microkernel

- Only the *essential* core OS functions should be in the kernel.
- Less essential services and applications are built on the microkernel and execute *in user mode*.
 - The operating system services are structured as a collection of independent processes
- Communication takes place between user modules using message passing.
- Example: Mach, QNX, L4 kernel, ...


Microkernels (2)

- Advantages
 - Extensibility, Modularity (Maintainability)
 - Small kernel makes it easy to debug and be more efficient
- Disadvantages: performance
 - Invoking services involve mode/process switches
 - Essential system services executing in user mode




Virtual Machine (1)

- In old days ...
 - An application program is bound to a specific platform which is 'ISA(instruction set architecture) + OS'



(a)



(b)

- What is a Virtual Machine (VM)?
 - Eliminates this real-platform constraint for high degree of portability and flexibility
 - Software for cross-platform compatibility


Virtual Machine (2)

- Why do we need VM?
 - Portability is essential in networked computing
 - Especially useful for mobile, wireless-download platforms where we can achieve consistent execution environment on diverse CPU/OS/hardware devices: e.g., Java VM, GVM, Brew, WIPI, ...
 - CPU innovations are often limited by old interfaces
 - New powerful CPUs that cannot run legacy x86 binaries are not viable on the market
 - Solution: run x86 binaries on high-performance, low-power CPUs
 - Single OS on a H/W may open a security hole
 - E.g., a server shared by different groups of users who want to be assured of a secure environment
 - Sandbox an OS that is not trusted, possibly because it is a system under development.
 - Virtual machines have other advantages for OS development, including better debugging access and faster reboots


Virtual Machine (3)

- VM Solution
 - Implementing a layer of software (VM) for virtualization
 - Mapping a virtual **guest** system to a real **host** system
- Two types of VM
 - **Process VM**: virtualization of individual processes
 - E.g., running x86 applications on Alpha CPU
 - **System VM**: virtualization of complete systems
 - E.g., running Linux (and its applications) on Windows

Running
Windows(guest)
on
a Linux host



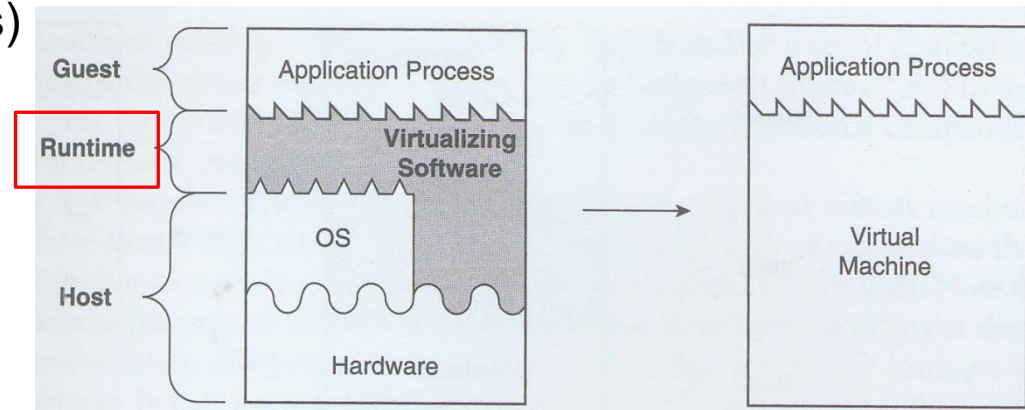
Running
Ubuntu
Linux(guest)
on a
Windows host



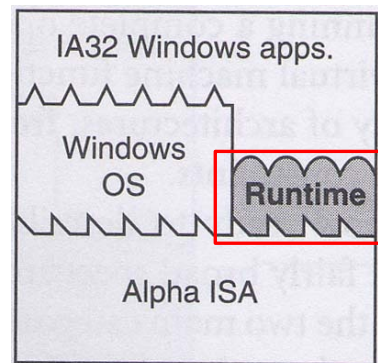

Virtual Machine (4)

• Process VM

- Run executables of different ISA or OS
- VM emulates ABI(application binary interface; user-level instructions & system calls)
- Called “runtime”



Java Virtual Machine:
Run Bytecode binaries on x86




Digital FX!32:
Run x86 binaries on Alpha

Virtual Machine (5)

• System VM


- Run whole OS(es) & executables
- VM emulates whole ISA (user-level & system-level)
- Called “VMM” or “Hypervisor”
(Virtual Machine Manager)



(a) One ISA is emulated by the other


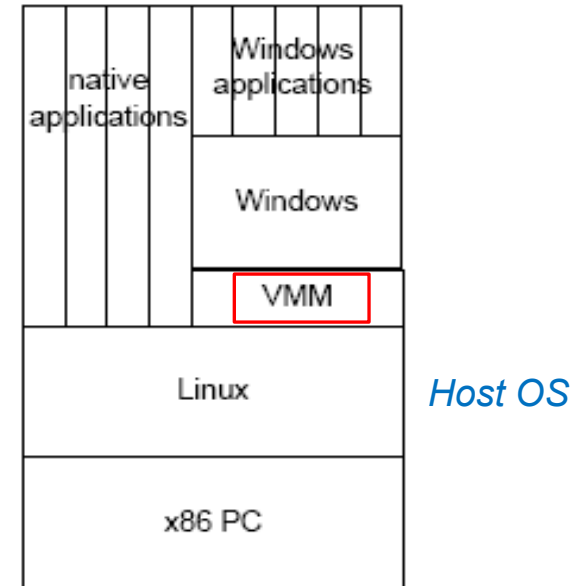
(b) Multiple platforms replicated on a single platform

(c) Multiple different platforms on a single platform

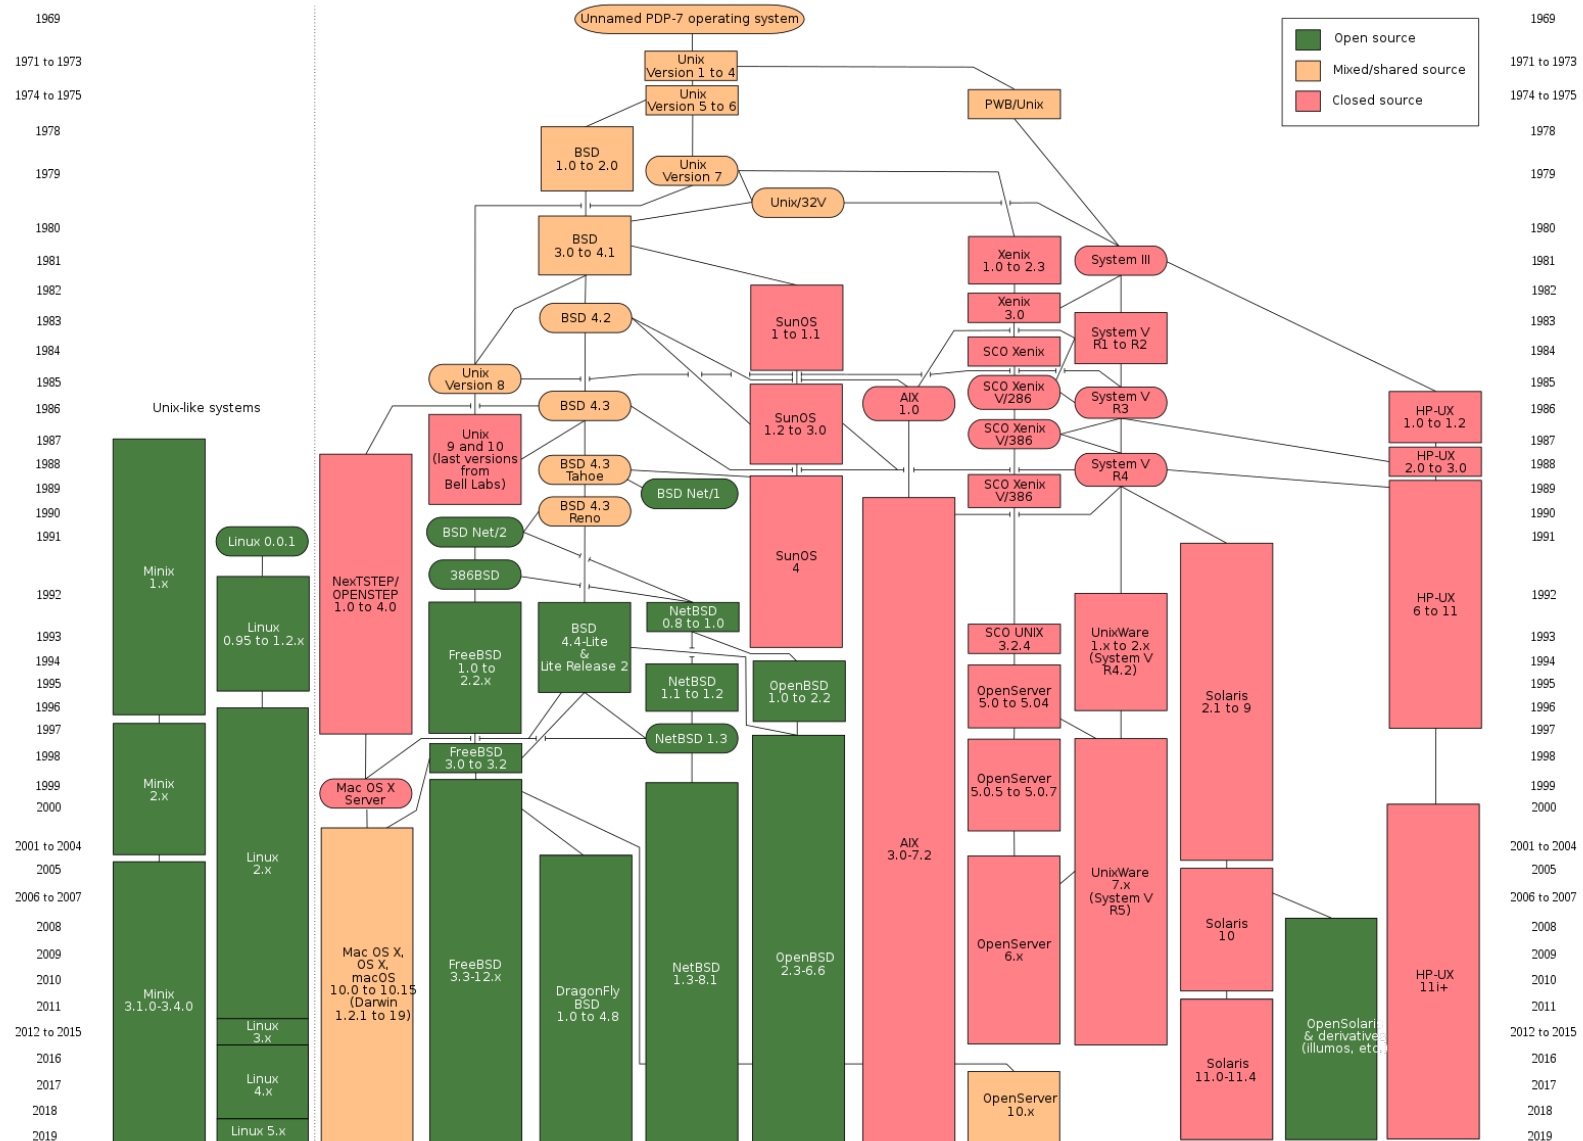


Virtual Machine (6)

- **System VM: Implementation issues**
 - Hosted or Stand-alone
- Hosted ('hosted hypervisor')
 - Runs as a process on an existing host OS
 - Rely on host OS for H/W interaction
 - VMWare Workstation, User-Mode Linux, Oracle VirtualBox, ...
- Stand-alone ('bare-metal/native hypervisor')
 - VMM on top of bare hardware
 - All H/W interactions done by VMM itself
 - Highly efficient
 - VMWare ESX, IBM z/VM, ...



History of UNIX (Wikipedia)



Linux

- UNIX-like OS
 - Linux is a clone of the operating system UNIX, written from scratch by Linus Torvalds with assistance from a loosely-knit team of programmers across the Net.
- A true UNIX kernel
 - But not a full UNIX operating system as it does not include all the UNIX applications such as compilers, windowing systems, ...
 - However, most of these programs are freely available!
- Source code under the GNU Public License (GPL)
 - Open and available to anyone to study.
 - <http://www.kernel.org/>

Some History

- GNU: Open Source before Linux
 - The concept of `free' software
 - Through '70 Richard Stallman advocates free software.
 - `free' as in freedom: i.e. free to use, distribute and modify.
 - 1984: Richard Stallman founded GNU
 - Goal: to produce free software.
 - GPL: ensure software freedom by copyright terms.
 - GNU software: UNIX-like program, **but no kernel!**
- Linux right on the start
 - From early start, Linus asked for volunteers on the Internet to help him develop Linux.
 - From the start, the source code has been freely available on the Internet.
 - GNU has lots of user-space programs but no kernel.

Why Linux? (1)

- Free and Open
 - Source code under the “GNU Public License”
 - Everyone has the right to use, copy and modify the programs free of charge.
 - C.f., Microsoft Windows, VxWorks, Solaris, ...
- Portable from mainframe to handhelds
 - Runs on low-end, cheap hardware platforms
 - Originally developed only for the PC; now runs on Alpha/Sparc/...
 - Minimum system requirements
 - CPU/RAM/Disk: 386+/8 MB/256 MB
- Compatibility and Efficiency
 - Highly compatible with many common OS
 - Powerful: the main Linux target is *efficiency*.
 - Small and compact kernel

Why Linux? (2)

- **Maturity**
 - ~30-year-old UNIX-type OS
 - Linux: since November 1991
 - High standard for source code quality
 - Stable, low failure rate and system maintenance time
- **Well supported**
 - Believe it or not, all commercial UNIX variants run on a restricted subset of hardware components

Key Features of Linux (1)

- Monolithic kernel
 - A large, complex program, composed of several logically different components
 - **High performance**: low message passing overhead
 - C.f., microkernel approach (modular approach) : CMU's Mach
- Supporting “modules”
 - To dynamically load and unload some portions of the kernel code on demand (typically, device drivers)
 - C.f, traditional Unix kernels: compiled and linked statically
 - Only SVR4.2 kernel has a similar feature.

Key Features of Linux (2)

- Kernel threading
 - “Kernel thread”: an execution context that can be independently scheduled on a common address space
 - Linux uses kernel threads in a very limited way to execute a few kernel functions
- Multithreaded application support
 - Linux defines its own version of **lightweight process**, which is different from those of (kernel thread based) Solaris or SVR4.
 - Lightweight process is the basic execution context.
 - Lightweight process is handled via the nonstandard **clone()** system call: “copy-on-write”

Key Features of Linux (3)

- Preemptive kernel
 - Starting from 2.6
 - Interleave execution flows while they are in privileged mode.
- Multiprocessor support
 - From Linux 2.2, yet make optimal use of SMP
- File system
 - The standard file system lacks some advanced features such as journaling, but more advanced file systems are available for Linux.

Linux Kernel Architecture

User Level



System Call Interface

Filesystem Manager

- 1.Ext2fs
- 2.proc
- 3.nfs

Memory Manager

Process Manager

- 1.Task Management
- 2.Scheduler
- 3.Signaling

Kernel Level

Buffer cache

Device Manager

- 1.block
- 2.character

Network Manager

- 1.Ipv6
- 2.ethernet

Device Interface

HW Level



Linux Kernel Source Tree

